

Parallel Segmented Merge and Its Applications to Two Sparse Matrix Kernels

Weifeng Liu, Norwegian University of Science and Technology, Norway

Hao Wang, Ohio State University, USA

Brian Vinter, University of Copenhagen



June 8th, 2018, Bergen, Norway

Background

- 1. Merge (sort) and its use**
- 2. A definition of segmented merge**
- 3. Merge in parallel**

Merge (sort)

- We consider to merge two sorted key-value ascending/descending sequences into one. A serial algorithm first creates two pointers running along the two sequences, compares the entries pointed by them, and saves the smaller/larger one to the resulting sequence.

Two sorted sequences

0	1	0	2
2	3	b	c

One sorted resulting sequence

0	0	1	2
2	b	3	c

Applying the left merging to sparse vector-vector addition:

$$\begin{array}{c} \boxed{2} \quad \boxed{3} \quad \boxed{} \quad \boxed{} \\ + \quad \boxed{b} \quad \boxed{} \quad \boxed{c} \quad \boxed{} \\ = \quad \boxed{2+b} \quad \boxed{3} \quad \boxed{c} \quad \boxed{} \end{array}$$

Use case 1: sparse matrix-matrix addition

- Add a sparse matrix A by a sparse matrix B , and obtain a resulting sparse matrix C .

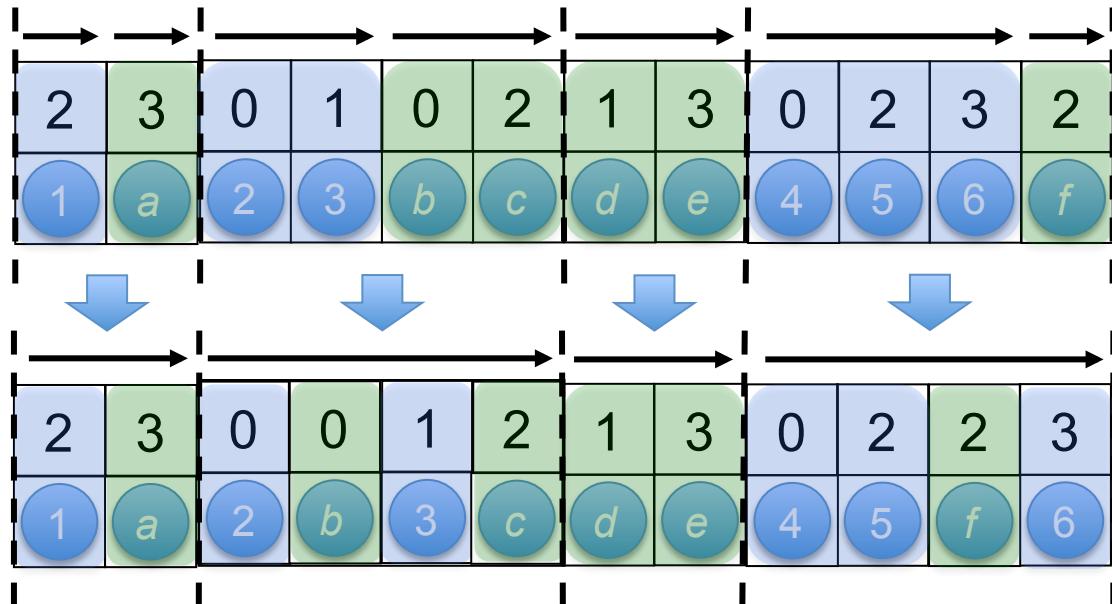
$$\begin{array}{c} \begin{matrix} & & & \\ & & & \\ & & 1 & \\ & & & \\ \hline & 2 & 3 & \\ & & & \\ \hline & 4 & & 5 & 6 \end{matrix} + \begin{matrix} & & & \\ & & & \\ & b & & a \\ & & c & \\ \hline & d & & e \\ & & f & \\ \hline \end{matrix} = \begin{matrix} & & & \\ & & & \\ & 1 & & a \\ & & & \\ \hline & 2+b & 3 & c \\ & & d & e \\ \hline & 4 & & 5+f & 6 \end{matrix} \end{array}$$

A
 (4×4)
sparse, 6 nonzeros B
 (4×4)
sparse, 6 nonzeros C
 (4×4)
sparse, 10 nonzeros

Use case 1: sparse matrix-matrix addition

- Each resulting row is merged from two rows.

$$\begin{array}{|c|c|c|} \hline & & 1 \\ \hline 2 & 3 & \\ \hline 4 & 5 & 6 \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline & & a \\ \hline b & c & \\ \hline d & e & \\ \hline f & & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & 1 & a \\ \hline 2+b & 3 & c & \\ \hline 4 & 5+f & 6 & \\ \hline \end{array}$$



Use case 2: sparse matrix-matrix multiplication

- Multiply a sparse matrix A by a sparse matrix B , and obtain a resulting sparse matrix C .

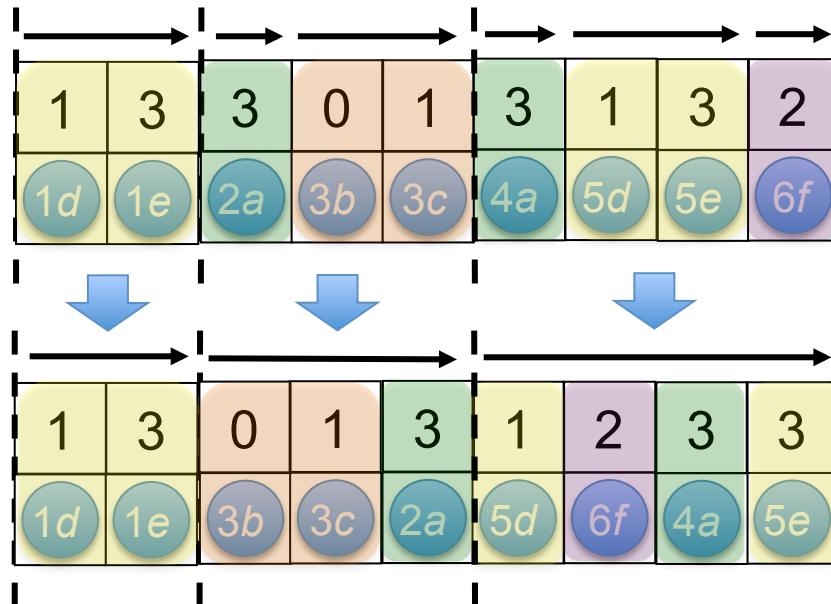
$$\begin{array}{c} \begin{array}{|c|c|c|c|} \hline & & 1 & \\ \hline 2 & 3 & & \\ \hline & & & \\ \hline 4 & & 5 & 6 \\ \hline \end{array} \times \begin{array}{|c|c|c|c|} \hline & & a & \\ \hline b & & c & \\ \hline & d & & e \\ \hline & & f & \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline & 1d & & 1e \\ \hline 3b & & 3c & 2a \\ \hline & 5d & 6f & 4a+5e \\ \hline \end{array} \end{array}$$

A
(4x4)
sparse, 6 nonzeros B
(4x4)
sparse, 6 nonzeros C
(4x4)
sparse, 8 nonzeros

Use case 2: sparse matrix-matrix multiplication

- Each resulting row is merged from multiple rows.

$$\begin{array}{|c|c|c|} \hline & & 1 \\ \hline 2 & 3 & \\ \hline 4 & 5 & 6 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline & & a \\ \hline b & c & \\ \hline d & e & f \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & 1d & 1e \\ \hline 3b & 3c & 2a \\ \hline 5d & 6f & 4a+5e \\ \hline \end{array}$$

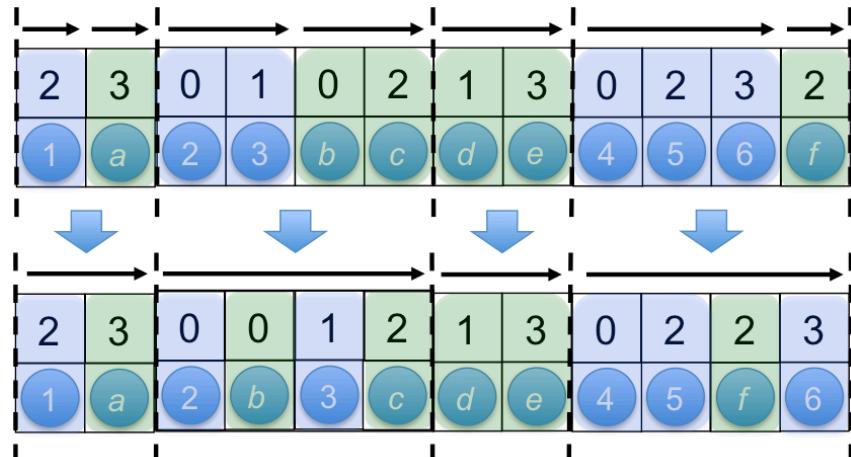


Background

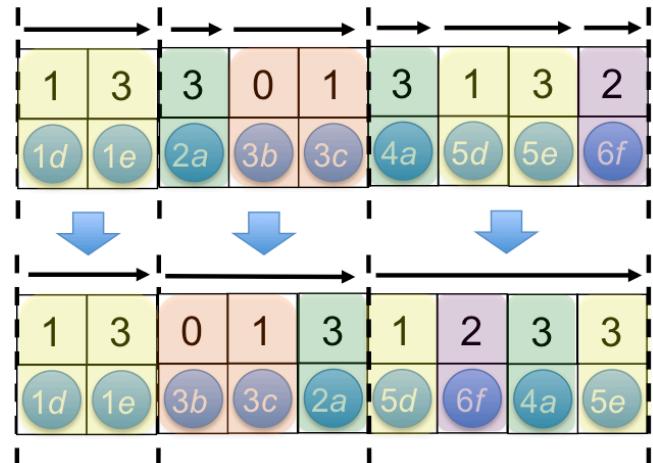
- 1. Merge (sort) and its use**
- 2. A definition of segmented merge**
- 3. Merge in parallel**

A definition of segmented merge (segmerge)

- Segmented merge operation merges q sub-segments into p segments.



$$q = 7, p = 4$$



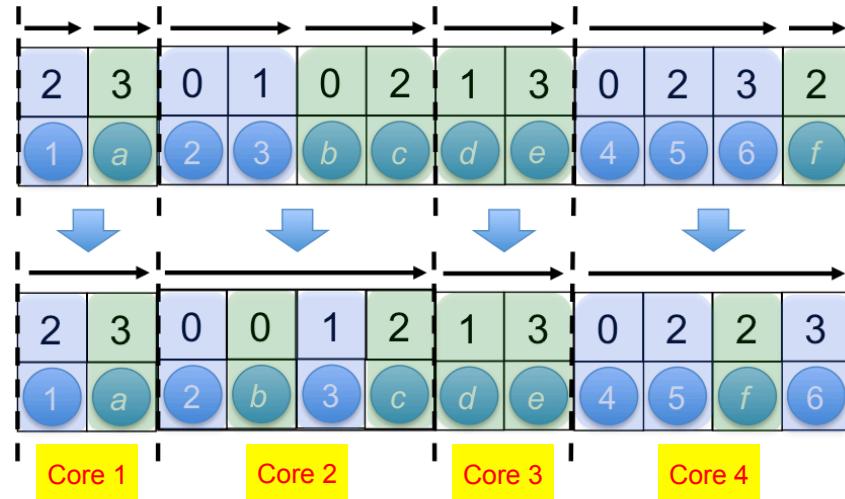
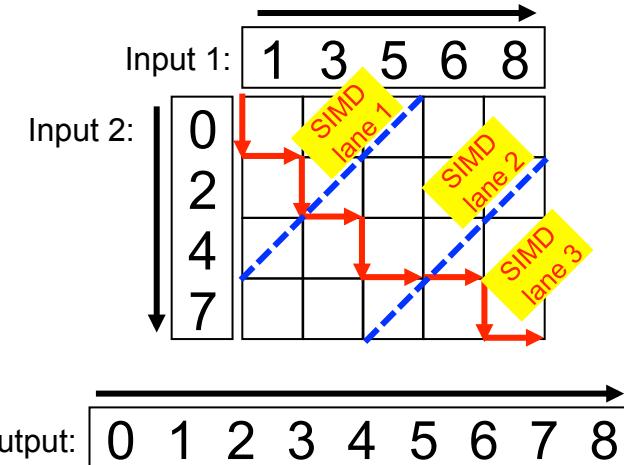
$$q = 6, p = 3$$

Background

1. Merge (sort) and its use
2. A definition of segmented merge
3. Merge in parallel

Merge in parallel

- There are two parallel merge methods:
 - (1) **merging two sub-segments**: merge-path works in a vectorized and balanced way,
 - (2) **merging multiple segments**: assigning each segment to a core



Though each segment can be merged in a balanced way, cores may have imbalanced workload due to very diverse segment sizes.

This motivates our segmerge algorithm.

Green, McColl, Bader. GPU Merge Path: a GPU Merging Algorithm. ICS '12.

Parallel Segmented Merge

1. Algorithm description
2. Experiments: SpTRANS and SpGEMM
3. Micro-benchmarks

Overview of our segmerge

Step 1. create level-1 and level-2 pointers

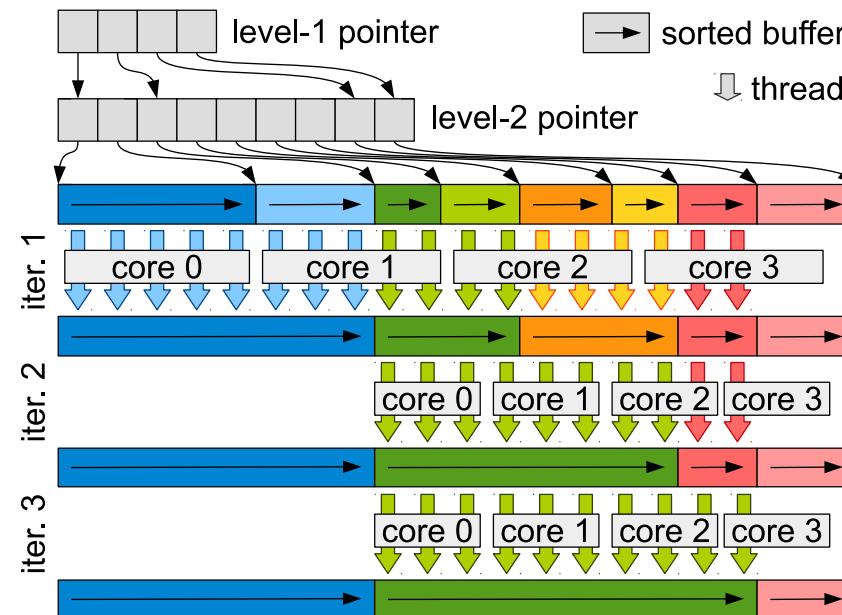
Step 2. fix each thread's workload (#entries) and compute #threads

Step 3. sub-segments scatter info to threads

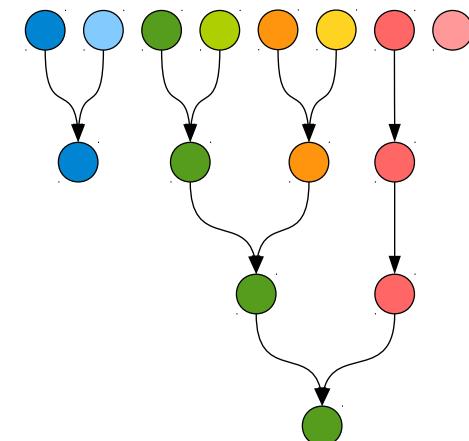
Step 4. threads gather info to find workload (similar to merge path)

Step 5. threads are evenly assigned to cores

Step 6. iterate steps 2- 5 until finish



Merge tree of the left segments



Parallel Segmented Merge

1. Algorithm description
2. Experiments: SpTRANS and SpGEMM
3. Micro-benchmarks

Experiment 1: sparse matrix transposition

- Transpose a sparse matrix A (in CSR) to a sparse matrix B (i.e., A^T , in CSR). This is equivalent to a conversion from CSR to CSC, or vice versa.

row pointer =

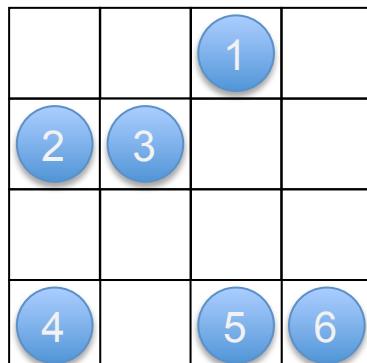
0	1	3	3	6
---	---	---	---	---

column index =

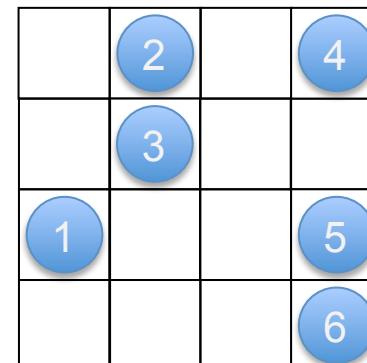
2	0	1	0	2	3
---	---	---	---	---	---

value =

1	2	3	4	5	6
---	---	---	---	---	---



->



row pointer =

0	2	3	5	6
---	---	---	---	---

column index =

1	3	1	0	3	3
---	---	---	---	---	---

value =

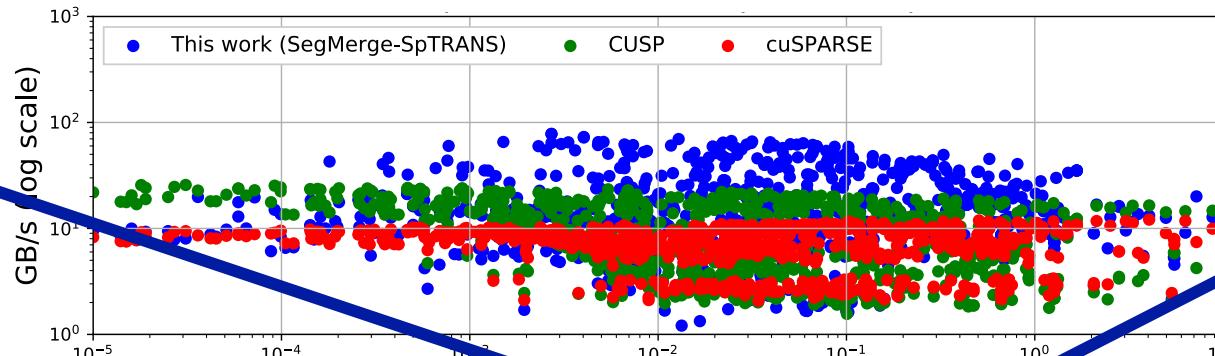
2	4	3	1	5	6
---	---	---	---	---	---



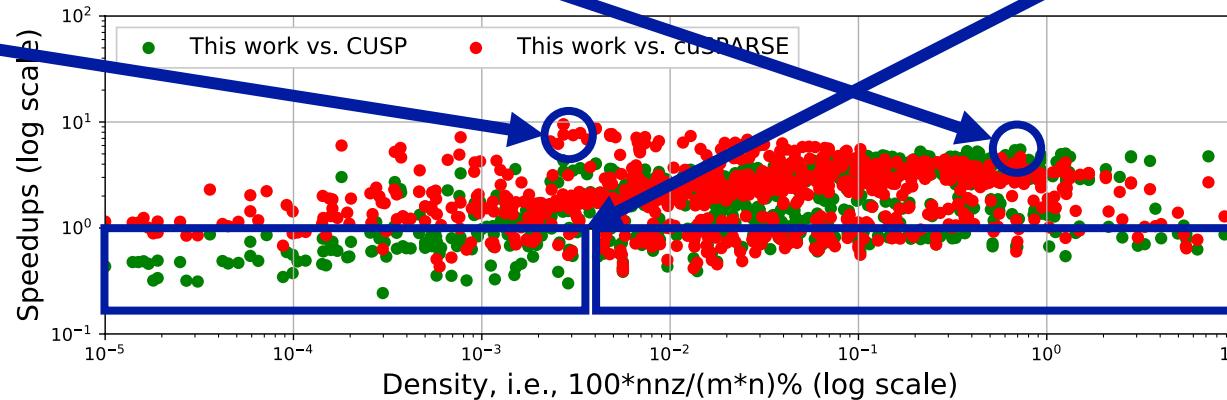
First collect entries,
then segmerge them.

Experiment 1: sparse matrix transposition

~7x over
CUSP



~10x over
cuSPARSE

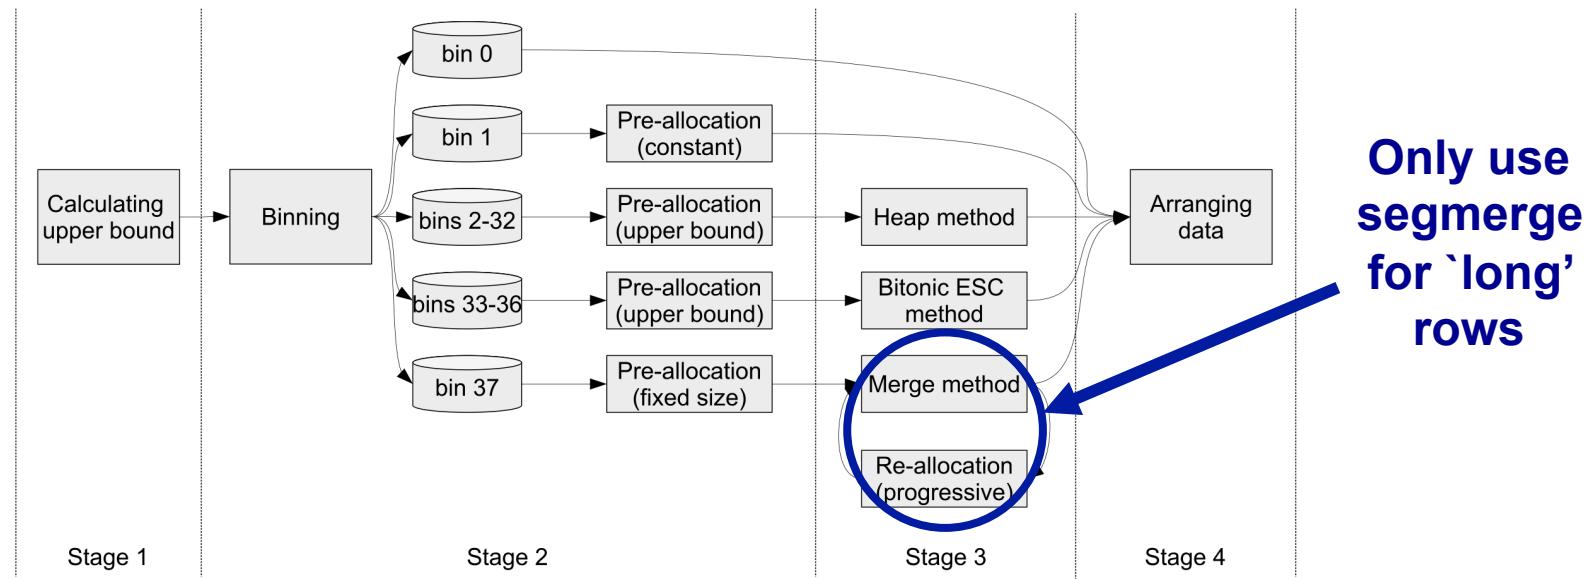


CUSP behaves
better when
very sparse

In some cases
slower than
CUSP and
cuSPARSE

GPU: Titan X
#matrices: 956

Experiment 2: sparse matrix-matrix multiplication



Weifeng Liu, Brian Vinter. **A Framework for General Sparse Matrix-Matrix Multiplication on GPUs and Heterogeneous Processors.** *JPDC*. 2015. (extended from *IPDPS14*)

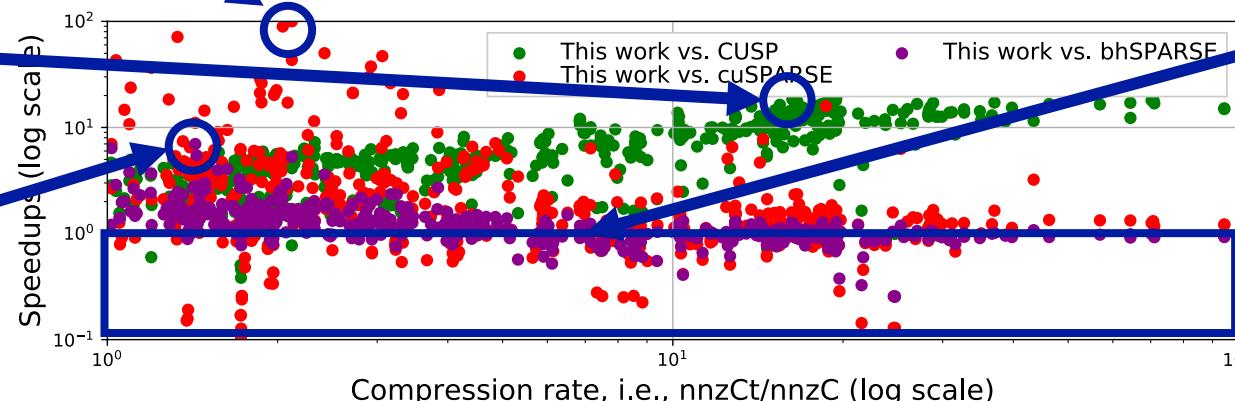
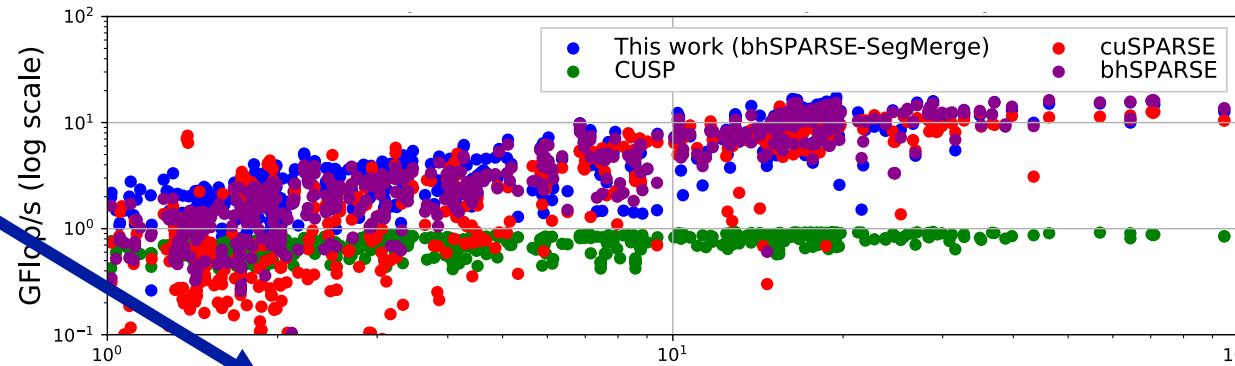
Weifeng Liu, Brian Vinter. **An Efficient GPU General Sparse Matrix-Matrix Multiplication for Irregular Data.** *IPDPS14*.

Experiment 2: sparse matrix-matrix multiplication

~100x over
cuSPARSE

~20x over
CUSP

~7x over
bhSPARSE



In some cases,
cuSPARSE
is faster, and
bhSPARSE
rarely gives
higher perf.

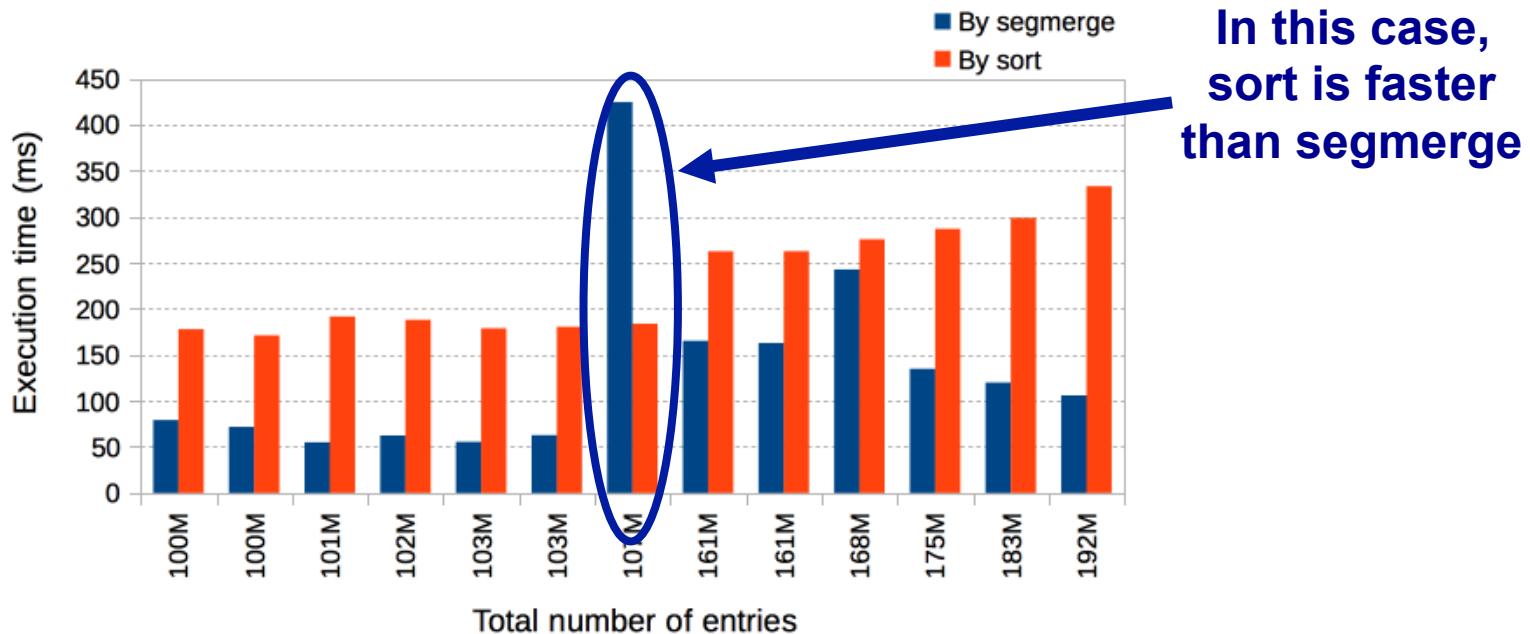
GPU: Titan X
#matrices: 956

Parallel Segmented Merge

1. Algorithm description
2. Experiments: SpTRANS and SpGEMM
3. Micro-benchmarks

Micro-benchmark 1: segmerge vs sort

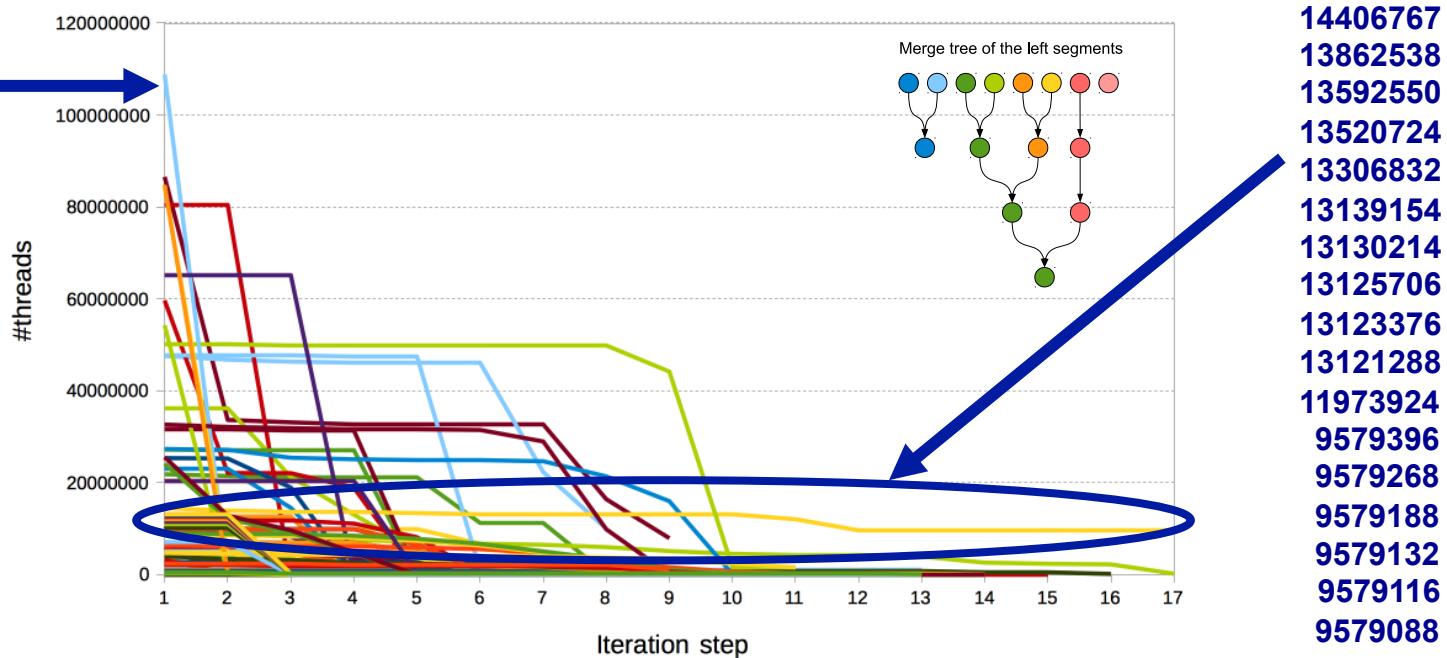
- In the SpGEMM test, if we replace segmerge with CUDA thrust-sort:



Micro-benchmark 2: #threads in iterations

- In the SpGEMM test, we record #threads issued in each iteration step:

108813822
552638
547690
515739
349187
32396
12215
6863
5987
5971
5963
5959
5956



Conclusion

- Though some parallel segmented operations (segsum[Blelloch TR93, Liu PARCO15], segsort [Hou ICS17], etc.) exist, parallel segmerge has been largely ignored.
- Our segmerge algorithm always evenly distributes entries to threads, thus runs in a very balanced way on GPU.
- SpTRANS and SpGEMM are largely accelerated by using segmerge at proper stages.
- According to micro-benchmarks, there are still some optimizations can be done.

T k u !

0	4	8	9
---	---	---	---

A y Q s n s ?

0	2	4	7	11	12	13
---	---	---	---	----	----	----